

Assignment 2

Austin Chau

HW: POS Tagging

John saw the saw and decided to take it to the table.

NNP VBD DT NN CC VBD TO VB PRP IN DT NN

Intro: Data

- ❑ Part of the Penn Treebank POS data set
 - ❑ Collection of sentences pretagged with POS
- ❑ Read the README file
 - ❑ `./train_hmm.py ptb.2-21.tgs ptb.2-21.txt > my.hmm`
- ❑ `ptb.2-21.tgs` → all the POS tags
- ❑ `ptb.2-21.txt` → all the sentences
- ❑ `train_hmm.py` → zip the two files together

Intro: Train_hmm

- ❑ Either perl or python is fine → they do the same thing
- ❑ Zip the POS and sentences together
- ❑ Generates my.hmm
 - ❑ List of emissions and transitions
 - ❑ Emissions → token and tag pair, and the probability assigned to it
 - ❑ $\text{float}(\text{emissions}[\text{tag}][\text{token}]) / \text{emissionsTotal}[\text{tag}]$
 - ❑ Transitions → previous token and current token pair, and the probability assigned to it
 - ❑ $\text{transitions}[\text{prevtag}][\text{tag}] / \text{transitionsTotal}[\text{prevtag}]$

Task 1

- ❑ Train the model on subsets of the training data of different sizes
 - ❑ Divide the dataset into several subsets (based on index or randomly)
 - ❑ Get different performance by increasing the number of subsets
 - ❑ Plot the performance
- ❑ Resizing can be done in `train_hmm.py`, where you don't use all the pairs taken from the `tagFile` and `tokenFile`

Plot

❏ pip install matplotlib

```
# importing the required module
import matplotlib.pyplot as plt


# x axis values
x = [1,2,3]
# corresponding y axis values
y = [2,4,1]

# plotting the points
plt.plot(x, y)

# naming the x axis
plt.xlabel('x - axis')
# naming the y axis
plt.ylabel('y - axis')


# giving a title to my graph
plt.title('My first graph!')

# function to show the plot
plt.show()
```



Plot

■ Learning curve



Task 1

- ❑ Generate a learning curve
- ❑ Give your thoughts about getting more POS-tagged data and how it would affect your system?


Task 2

- ❑ Come up with a way to improve the model
- ❑ Build a replacement for the `train_hmm.{py,pl}` script (using HMM)
- ❑ Write your own Viterbi algorithm in Python (recommended)

Implement a trigram HMM

- ❑ Learn from `train_hmm.py`
- ❑ Keep the observations for each state
- ❑ Add two initial states before the first token
- ❑ Update previous two states in training

Viterbi for trigram



Viterbi for trigram

Input: a sentence $x_1 \dots x_n$, parameters $q(s|u, v)$ and $e(x|s)$.

Definitions: Define \mathcal{K} to be the set of possible tags. Define $\mathcal{K}_{-1} = \mathcal{K}_0 = \{*\}$, and $\mathcal{K}_k = \mathcal{K}$ for $k = 1 \dots n$.

Initialization: Set $\pi(0, *, *) = 1$.

Algorithm:

- For $k = 1 \dots n$,

- For $u \in \mathcal{K}_{k-1}, v \in \mathcal{K}_k$,

$$\pi(k, u, v) = \max_{w \in \mathcal{K}_{k-2}} (\pi(k-1, w, u) \times q(v|w, u) \times e(x_k|v))$$

$$bp(k, u, v) = \arg \max_{w \in \mathcal{K}_{k-2}} (\pi(k-1, w, u) \times q(v|w, u) \times e(x_k|v))$$

- Set $(y_{n-1}, y_n) = \arg \max_{u \in \mathcal{K}_{n-1}, v \in \mathcal{K}_n} (\pi(n, u, v) \times q(\text{STOP}|u, v))$
- For $k = (n-2) \dots 1$,

$$y_k = bp(k+2, y_{k+1}, y_{k+2})$$

- **Return** the tag sequence $y_1 \dots y_n$

Implement a trigram HMM

- ❑ Use log
 - ❑ Initialization
 - ❑ Observations out of vocabulary
 - ❑ Overflow
- ❑ Use dictionaries in Python

Smooth the probability estimates

- ❑ To solve data sparsity

$$P(t_i|t_{i-1}t_{i-2})$$

- ❑ Replace every tag by its first letter only, the tags would still be meaningful, only more coarse
- ❑ Estimate the probability by combining more robust but weaker estimators

$$P(t_i|t_{i-1}t_{i-2}) = \lambda_3 \hat{P}(t_i|t_{i-1}t_{i-2}) + \lambda_2 \hat{P}(t_i|t_{i-1}) + \lambda_1 \hat{P}(t_i)$$

How to choose λ ?

function DELETED-INTERPOLATION(*corpus*) **returns** $\lambda_1, \lambda_2, \lambda_3$

$\lambda_1 \leftarrow 0$

$\lambda_2 \leftarrow 0$

$\lambda_3 \leftarrow 0$

foreach trigram t_1, t_2, t_3 with $C(t_1, t_2, t_3) > 0$

depending on the maximum of the following three values

case $\frac{C(t_1, t_2, t_3) - 1}{C(t_1, t_2) - 1}$: increment λ_3 by $C(t_1, t_2, t_3)$

case $\frac{C(t_2, t_3) - 1}{C(t_2) - 1}$: increment λ_2 by $C(t_1, t_2, t_3)$

case $\frac{C(t_3) - 1}{N - 1}$: increment λ_1 by $C(t_1, t_2, t_3)$

end

end

normalize $\lambda_1, \lambda_2, \lambda_3$

return $\lambda_1, \lambda_2, \lambda_3$

Task 2

- ❑ Describe your approach clearly
- ❑ Provide the performance of your model on ptb.22.*
- ❑ Run your tagger on ptb.23.txt (the test data) and turn in the code and output
- ❑ Do not search ptb23.tag

Task 3: Different Languages

- ❑ Run the baseline model and your model from Task 2 on Japanese and Bulgarian
- ❑ Report and explain the results
 - ❑ Why are the baseline model and your model performing better/worse on Japanese and/or Bulgarian?

Bonus Task:

- ❑ Current:
 - ❑ We have only been reporting #errors compared to the test data in terms of words and sentence.
- ❑ Are there better metrics you can do to evaluate the difference between your model and the baseline?
 - ❑ What model can you use instead of just tallying up mismatch with the test set?

Bonus task

- ❑ In your answer
 - ❑ Your modified script
 - ❑ Your results in the PDF
 - ❑ Your explanation:
 - ❑ Why you get this results
 - ❑ Why your metric is better